

Contents

Deep Learning HW1	1
David Pitts 313371080	1
Aner Zakobar 203340302	1
Parctical part , Part 1 , Q1:	3
Part I	5
Question 2	5
Discussion SGD variants :	10
Part 1 Question 3 , Dataset "Peaks"	10
Learning rate and batch sizes experiments for peaks dataset	14
GMM dataset	15
Part 2	15
Jacobain test feed foward network	16
Part 2 Question 1	16
Part 2 Question 2	19
Jacobin test for ResNet	20
Part 2 , Question 3 , Gradient test feed forward	26
Part 2 Question 3 , Gradient test ResNet	28
Part 2 Question 4 , optimzation of feed forward nernal network on Peaks dataset	31
Experiments and failures network design :	35
Hyper parameter searching Peaks data :	35
Depth and epochs:	35
Learning rate :	36
LR=0.001	36
LR=0.1	36
Batch Size :	36
Batch-size == 8 :	36
Batch-size == 32	36
Question 2 Part 4 :	36
Stability and depth :	40
Depth = 2 network width =100 , iter =150	40
Depth =3 , network width=20 , iter =150	41
Depth =4 , network width=10 , iter =150	41
Question 2.5 - only 200 samples	41
Experiments	43
depth of 3 , width 20 , lr = 0.01 , batchsize =8	43
depth of 4 , width 10 , lr = 0.01 , batchsize = 8	43
depth of 2 , width 100 , lr = 0.01 , batchsize =8	43
How did the result changes	43

Deep Learning HW1

David Pitts 313371080

Aner Zakobar 203340302

```
1 import pandas as pd
import numpy as np
3 import matplotlib.pyplot as plt
```

```
%matplotlib inline
5 import scipy.sparse
```

```
7 import numpy as np
```

Gradient checking

```
1 def softmax(z):
    z -= np.max(z)
    3 sm = (np.exp(z).T / np.sum(np.exp(z),axis=1)).T
    return sm
```

```
y_hot=pd.get_dummies(np.arange(10))
```

```
2
def getLoss(w,x,y):
    4 """
        Forward for a simple softmax regression ,
    6 :param w: weights
        :param x:
    8 :param y:
        :return: the loss and the gradient .
    10 """
    y_mat=np.array(y_hot[y])
    12 scores = np.dot(x,w)
    prob = softmax(scores)
    14 loss = (-1) * np.sum(y_mat.T * np.log(prob))
    grad = -(1) * (x.T@(y_mat.T - prob))
    16 return loss,grad
```

```
def gradient_test_softmax_reg():
    2 """
        Performs a gradient test for softmax regression ,
    4 the data and weights are intilized randomly inside the function.
        :return:
    6 """
    values_linear=[]
    8 values_quard=[]
    epsilon=2

    10
    w=np.random.random((10,10))
    12 x=np.random.random((1,10))
    y=np.random.randint(0,10,size=(1))

    14
    f0 = getLoss(w, x, y)[0]
    16 grad = getLoss(w, x, y)[1]

    18 b = np.random.random(w.shape)
    b= b/np.linalg.norm(b,ord=2)

    20
    gradient_test_convergence(b, epsilon, f0, grad, values_linear, values_quard, w, x, y)

    22
    24 values_quard=np.array(values_quard)
    values_linear=np.array(values_linear)
    26 return values_quard,values_linear

    28
def gradient_test_convergence(b, epsilon, f0, grad, values_linear, values_quard, w, x, y):
    30 """
        Performs the loop which saves the values of the gradient test for each iteration.
```

```

32 :param b: random vector of weights
33 :param epsilon:
34 :param f0: the value of the function with out epsilon
35 :param grad: the gradient of the function without epsilon
36 :param values_linear: list of the linear values
37 :param values_quad: list of the quad values
38 :param w: wieghts
39 :param x:
40 :param y:
41 :return:
42 """
43
44 for i in range(10):
45     eps = epsilon * (np.power(0.5, i))
46     curr_b = b * eps
47     w_eps = w + curr_b
48     grad_moved = (eps * b * grad).sum()
49     val_linear = np.abs(getLoss(w_eps, x, y)[0] - f0)
50     val_quad = getLoss(w_eps, x, y)[0] - getLoss(w, x, y)[0] - grad_moved
51     values_linear.append(val_linear)
52     values_quad.append(val_quad)

```

Parctical part , Part 1 , Q1:

In the tables below , the first colum demonstarte the values of the gradient test and the second row demonstrates the ratio between following iterations , and as we expected , the quad test results in a quadratic ratio while the linear test results in a square ratio

As we can see , the quad ratio is 4 while the linear ratio is 2 , which confirms that our gradient for softmax regression is correct.

```

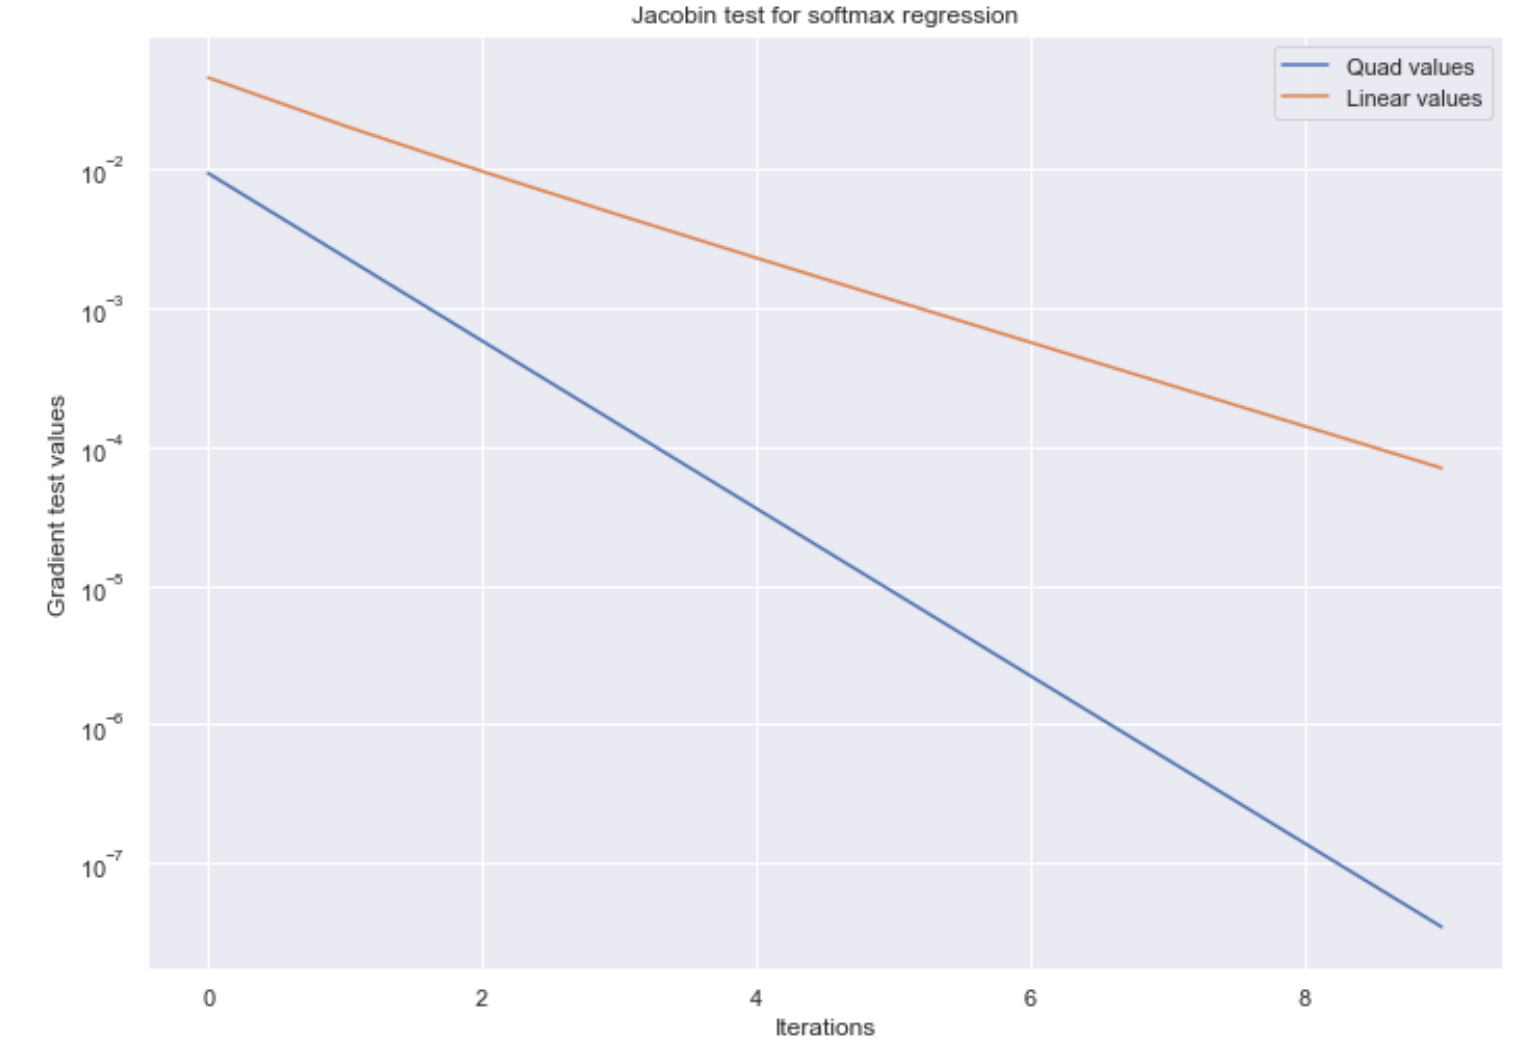
from IPython.core.display_functions import display
2
def gradient_test_demonstration():
3     """
4     Performs the gradient test and displays the data , the first column is the value at
5     each iteration and the second is the ration
6     between two following iterations.
7     :return:
8     """
9     values_quad, values_linear=gradient_test_softmax_reg()
10    quad_graident_test=pd.DataFrame(data=[values_quad, values_quad[:-1]/values_quad[1:]]).T
11    quad_graident_test.columns=["quadratic", "F(i)/F(i+1)"]
12    linear_gradient_test=pd.DataFrame(data=[values_linear, values_linear[:-1]/values_linear[1:]]).T
13    linear_gradient_test.columns=["linear", "F(i)/F(i+1)"]
14    display(quad_graident_test)
15    display(linear_gradient_test)
16    plt.plot(values_quad, label="Quad values")
17    plt.plot(values_linear, label="Linear values")
18
19    plt.legend()
20    plt.yscale("log")
21    plt.title("Jacobin test for softmax regression")
22    plt.ylabel("Gradient test values")
23    plt.xlabel("Iterations")
24    plt.show()
gradient_test_demonstration()

```

	quadratic	F(i)/F(i+1)
0	9.243172e-03	4.014270
1	2.302578e-03	4.008017
2	5.744932e-04	4.004230

5	3	1.434716e-04	4.002170
	4	3.584844e-05	4.001099
7	5	8.959650e-06	4.000553
	6	2.239603e-06	4.000277
9	7	5.598619e-07	4.000139
	8	1.399606e-07	4.000070
11	9	3.498955e-08	NaN

1		linear	F(i)/F(i+1)
	0	0.044999	2.229824
3	1	0.020181	2.121258
	2	0.009514	2.062335
5	3	0.004613	2.031610
	4	0.002271	2.015918
7	5	0.001126	2.007988
	6	0.000561	2.004001
9	7	0.000280	2.002002
	8	0.000140	2.001002
11	9	0.000070	NaN



Part I

Question 2

Write the code for minimizing an objective function using SGD or some other SGD variant (SGD with momentum, for example). Demonstrate and verify that your optimizer works on a small least squares example (add plots and submit the code itself).

We chose to implement 3 variation of SGD , the classic version , nestrov and mounemtoum as seen in the following section

```
1 def
  SGD(x,y,w,loss_func,batch_size=16,lr=1e-5,beta=0.9,itors=50000,nestrov=False,mounemtoum=False)
  """
3   An implementation of SGD with his variants , given data and weights he'll update the
   weight
   accordingly.
5   :param x:
6   :param y:
7   :param w:
8   :param loss_func:
9   :param batch_size:
10  :param lr:
11  :param beta:
12  :param iters:
13  :param nestrov:
14  :param mounemtoum:
15  :return:
   """
17  loss=[]
18  prev_grad=0
19  for i in range(itors):
20      if nestrov:
21          grad,val=loss_func(x,y,w-beta*prev_grad)
22      else:
23          grad,val=loss_func(x,y,w)
24
25      if nestrov or mounemtoum:
26          grad=beta*prev_grad+lr*grad
27          prev_grad=grad
28      else:
29          grad=lr*grad
30
31      w=w-grad
32      loss.append(val)
33  return w, loss
```

Small sample to verify that it works


```
1 import seaborn as sns
  sns.set(rc={'figure.figsize':(11.7,8.27)})
3
4 def loss_func_line_fitting(x,y,w):
5     """
6     giving noisy data and weights , we perform one iteration of their gradients and values
7     for simple
8     line fitting , i.e. we calculate the loss function and the gradient.
9     :param x:
10    :param y:
11    :param w:
12    :return:
   """
13    a=w[0]
```

```

15     b=w[1]
16     val=np.square(a*x+b-y).mean()
17     grad_a=(2*(a*x+b-y)*x).mean()
18     grad_b=2*(a*x+b-y).mean()
19     grad=np.array([grad_a,grad_b])
20     return grad,val
21 def simple_line_fitting(to_plot=True):
22     """
23     A simple line fitting which includes the init of random data and weights , followed by
24     their optimization using SGD with different variants.
25     :param to_plot:
26     :return:
27     """
28     w=np.random.random(2)
29     x=np.arange(1,50)
30     y=np.arange(1,50)+np.random.normal(size=49)*0.5
31     if to_plot==True:
32         sns.lineplot(x=x,y=y)
33         plt.xlabel("x")
34         plt.ylabel("y")
35         plt.title("Noisy data for line fitting , sampeled from  $y=x$  , guassain noise with
36                  $\sigma=0.5$ ")
37         plt.show()
38
39     w_n,loss_n=SGD(x,y,w,loss_func_line_fitting,lr=0.0001,nestrov=True)
40     w_m,loss_m=SGD(x,y,w,loss_func_line_fitting,lr=0.0001,mounemtoum=True)
41     w,loss=SGD(x,y,w,loss_func_line_fitting,lr=0.0001)
42
43     if to_plot==True:
44         a=w[0]
45         b=w[1]
46         print("SGD standart a = {} , b ={}".format(a,b))
47
48         plt.plot(loss[:100])
49         plt.title("Converges of the loss function")
50         plt.xlabel("Epochs")
51         plt.ylabel("Loss")
52         plt.yscale("log")
53         plt.show()
54         y_hat=a*x+b
55         plt.plot(x,y_hat,label=" $\hat{y}=\hat{a}x+\hat{b}$ ")
56         sns.lineplot(x=x,y=y,label="noisy sampel from  $y=x$ ")
57         sns.lineplot(x=x,y=np.arange(1,50),label=" $y=x$ ")
58         plt.xlabel("x")
59         plt.ylabel("y")
60
61     return w_n,w_m,w
62
63 1 simple_line_fitting()
64 print("")


```


Noisy data for line fitting , sampeled from $x = y$, guassain noise with $\sigma = 0.5$



SGD standart a = 1.0067496501143596 , b = -0.17453831850205515

Converges of the loss function





Comparing SGD variants for the least square example, because we know the optimal parameters of a and b ($a=1, b=0$) it's easy to assess the competitiveness of the SGD variation to the current problem.

```
def compare_SGD_variants_line_fitting():
    """
    Comparing the learned weights of  $a$  and  $b$  to their optimal values ( $a=1, b=0$ ), by fitting
    a simple noisy line 30 times, 3 times for each variation.
    :return:
    """
    w_l=[]
    w_lm=[]
    w_ln=[]

    for i in range(10):
        #receive weights for each simple line fitting and save them
        w_n,w_m,w=simple_line_fitting(to_plot=False)
        w_l.append(w)
        w_lm.append(w_m)
        w_ln.append(w_n)

    wl=np.array(w_l)
    w_lm=np.array(w_lm)
    w_ln=np.array(w_ln)

    df_a=pd.DataFrame([wl[:,0],w_lm[:,0],w_ln[:,0]]).T
    df_a.columns=["a_hat SGD","a_hat SGD momentum","a_hat SGD nesterov"]
```

```

24 df_b=pd.DataFrame([wl[:,1],w_lm[:,1],w_ln[:,1]]).T
25 df_b.columns=["b_hat SGD","b_hat SGD mounentuom"," b_hat SGD nesterov"]
26
27 display(df_a)
28 display(df_b)
29
30 comapre_SGD_variants_line_fitting()

```

```

1  a_hat SGD  a_hat SGD mounentuom  a_hat SGD nesterov
0  1.000537      1.001957      1.001957
3 1  0.996107      0.998472      0.998472
2  0.996033      0.996625      0.996625
5 3  0.997485      0.997656      0.997656
4  1.002971      1.003958      1.003958
7 5  0.994313      0.995146      0.995146
6  0.995111      0.997015      0.997015
9 7  0.997278      0.999292      0.999292
8  1.002259      1.003421      1.003421
11 9  0.996964      0.997293      0.997293

```

```

1  b_hat SGD  b_hat SGD mounentuom  b_hat SGD nesterov
0  0.017465      -0.029383      -0.029383
3 1  0.175874      0.097835      0.097835
2  -0.032174      -0.051706      -0.051706
5 3  0.146998      0.141375      0.141375
4  -0.027643      -0.060202      -0.060202
7 5  0.103172      0.075677      0.075677
6  0.112265      0.049452      0.049452
9 7  0.193069      0.126639      0.126639
8  -0.123447      -0.161764      -0.161764
11 9  0.265536      0.254691      0.254691

```

Discussion SGD variants :

Each have their strengths and weaknesses , it seems that for the simple problem of line fitting there isn't much different at the converging values for SGD with momentum or nesterov variant , but they do differ from SGD , managing to outperform or vice versa.

```

1 import scipy.io
mat = scipy.io.loadmat('PeaksData.mat')

```

Part 1 Question 3 , Dataset "Peaks"

```

1 y_train=mat["Ct"]
  x_train=mat["Yt"]
3
4 y_val=mat["Cv"]
5 x_val=mat["Yv"]
6
7 # Adding bias since its regression
  x_train=np.vstack([x_train,np.ones(25000)])
9 x_val=np.vstack([x_val,np.ones(6250)])

```

```

def getLoss1(x,y,w,val=False):
    """
    Loss for softmax regression with an option to return the predicted label
    :param x:

```

```

6      :param y:
      :param w:
      :param val:
8      :return:
      """
10     x=x.T
11     m = x.shape[0]
12     y_mat=y.T
13     scores = np.dot(x,w)
14     prob = softmax(scores)
15     loss = (-1 /m) * np.sum(np.sum(y_mat * np.log(prob)))
16     grad = (-1/m)*np.dot(x.T,(y_mat - prob))
17     if val==True:
18         return np.argmax(prob,axis=1)
19     return grad,loss
20
21
22
23 def
24 SGD(x,y,w,loss_func,batch_size=16,lr=1e-5,beta=0.9,itors=50000,nestrov=False,mounemtoum=False)
25     """
26     Same implementation as SGD as before , but supports batch-sizes now by picking random
27     indices in the
28     dataset length.
29     :param x:
30     :param y:
31     :param w:
32     :param loss_func:
33     :param batch_size:
34     :param lr:
35     :param beta: beta hyper parameter for mounemtoum
36     :param iters:
37     :param nestrov: flag
38     :param mounemtoum: flag
39     :return:
40     """
41     loss=[]
42     prev_grad=0
43     for i in range(itors):
44         #Batch sampling
45         indices=np.random.randint(0,25000,batch_size)
46         if nestrov:
47             grad,val=loss_func(x[:,indices],y[:,indices],w-beta*prev_grad)
48         else:
49             grad,val=loss_func(x[:,indices],y[:,indices],w)
50         if nestrov or mounemtoum:
51             grad=beta*prev_grad+lr*grad
52             prev_grad=grad
53         else:
54             grad=lr*grad
55         w=w-grad
56         loss.append(val)
57     return w, loss
58
59 def
60 optimize_f2(x,y,loss_func,lr=0.001,nestrov=False,moumentoum=False,plot=True,batch_size=32,data

```

```

62 """
    Performs an optimization of a given dataset using softmax regression , supports Peaks
        and GMM only.
    :param x:
    :param y:
    :param loss_func:
    :param lr:
    :param nestrov: flag for SGD variant
    :param moumentoum: flag for SGD varaint
    :param plot: flag which indicate plotting
    :param batch_size:
    :param data_set_peaks: flag to indicate which dataset
    :return:
    """
74 if data_set_peaks:
76     w=np.zeros([3,5])
    else:
78     w=np.zeros([6,5])
    loss=[]
80     acc_train=[]
    acc_val=[]
82     for i in range(200):
        lr=lr*np.power(0.9,i)
84
86     w,loss_curr=SGD(x,y,w,loss_func,lr=lr,itters=25000//batch_size,batch_size=batch_size,nestrov
88
89     samples=25000
    agreements = getLoss1(x, y, w, val=True) == np.argmax(y, axis=0)
90     loss.append(np.array(getLoss1(x, y, w)[1]))
92
93     acc_train.append(agreements.sum()/samples)
94
95     samples=6250
96
97     agreements = getLoss1(x_val, y_val ,w, val=True) == np.argmax(y_val, axis=0)
    acc_val.append(agreements.sum()/samples)
98
99 if plot :
100     plt.plot(acc_val,label="acc validation")
    plt.plot(acc_train,label="acc train")
102     plt.legend()
    plt.title("Accuracy on Peaks data set , using SGD nestrov")
104     plt.xlabel("Epochs")
    plt.ylabel("Accuracy")
106     plt.show()
108
109     plt.title("function loss on Peaks data set , using SGD nestrov")
    plt.xlabel("Epochs")
110     plt.ylabel("loss")
    plt.plot(loss)
112     plt.show()
114
    return w, loss , acc_val,acc_train

def comapre_sgd_variants(data_set_peaks=True):
2    """
    Comapres different SGD varaints on a given dataset trained with softmax regression

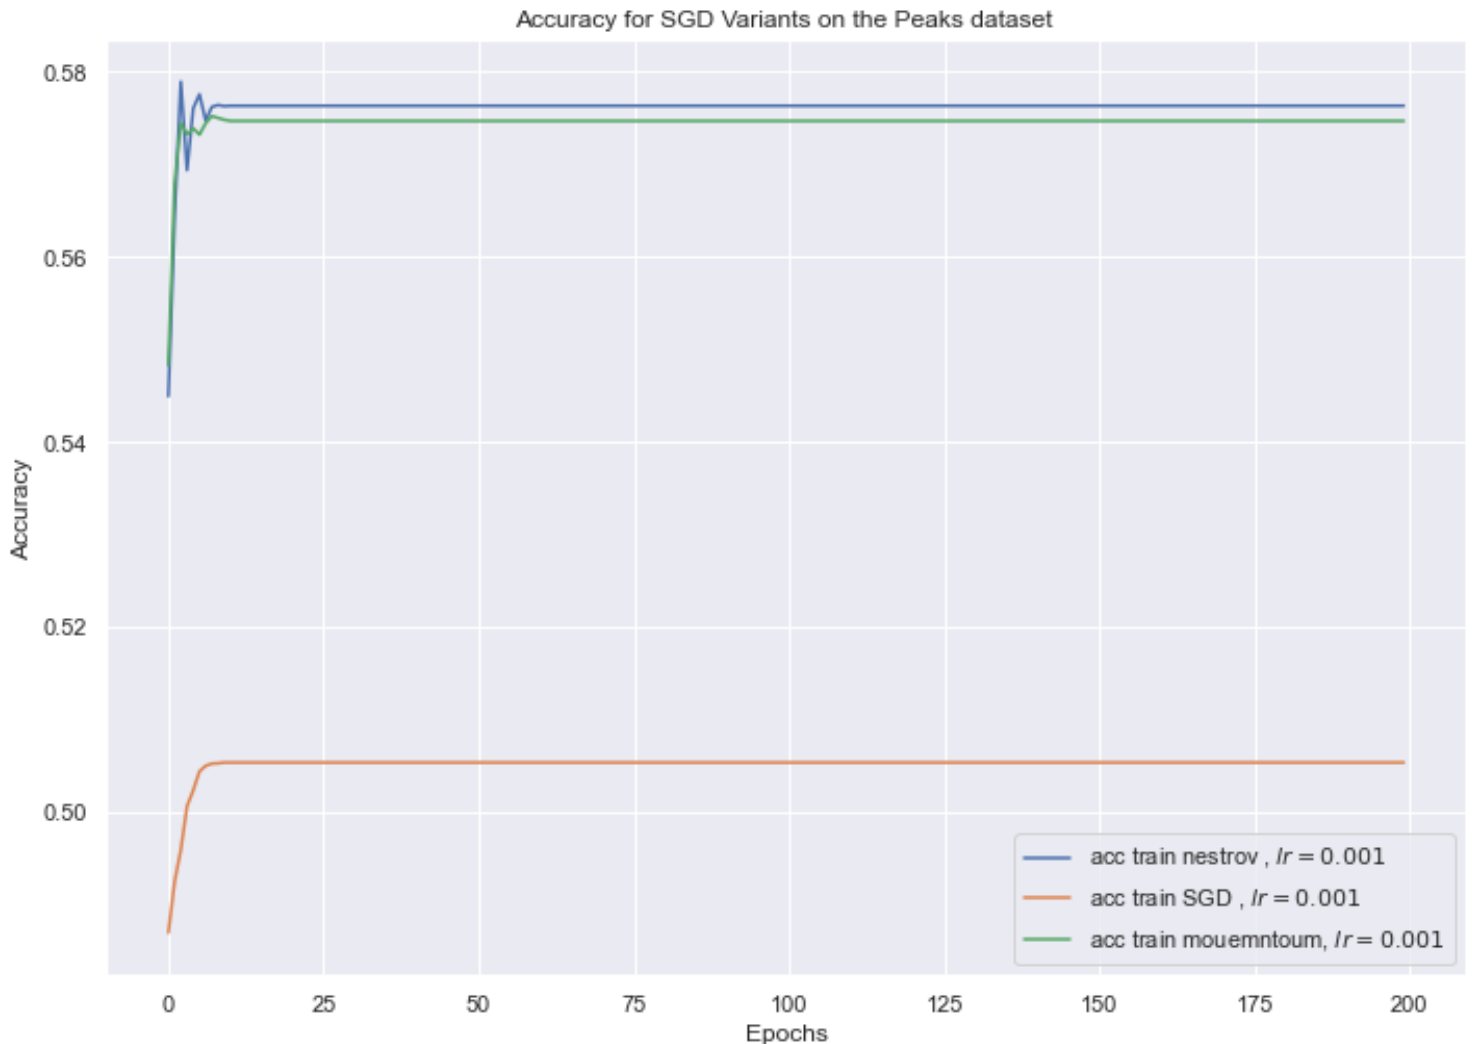
```

```

4  :param data_set_peaks:
   :return:
6  """
   __,val_sgd,train_sgd=optimize_f2(x_train,y_train,getLoss1,plot=False,data_set_peaks=data_set_peaks)
8  __,val_v,train_v=optimize_f2(x_train,y_train,getLoss1,nestrov=True,plot=False,data_set_peaks=data_set_peaks)
10 __,val_m,train_m=optimize_f2(x_train,y_train,getLoss1,mouementoum=True,plot=False,data_set_peaks=data_set_peaks)

12 if data_set_peaks:
13     data_set="Peaks"
14 else:
15     data_set="GMM"
16 plt.plot(train_v,label="acc train nestrov , $lr=0.001$")
17 plt.plot(train_sgd,label="acc train SGD , $lr=0.001$")
18 plt.plot(train_m,label="acc train mouemntoum, $lr=0.001$")
19 plt.title(f"Accuracy for SGD Variants on the {data_set} dataset ")
20 plt.xlabel("Epochs")
21 plt.ylabel("Accuracy")
22 plt.legend()
23 plt.show()
24 comapre_sgd_variants()

```



Preperation of the data for the GMM dataset , training using softmax regression.

```
mat=scipy.io.loadmat("GMMData.mat")
```

```

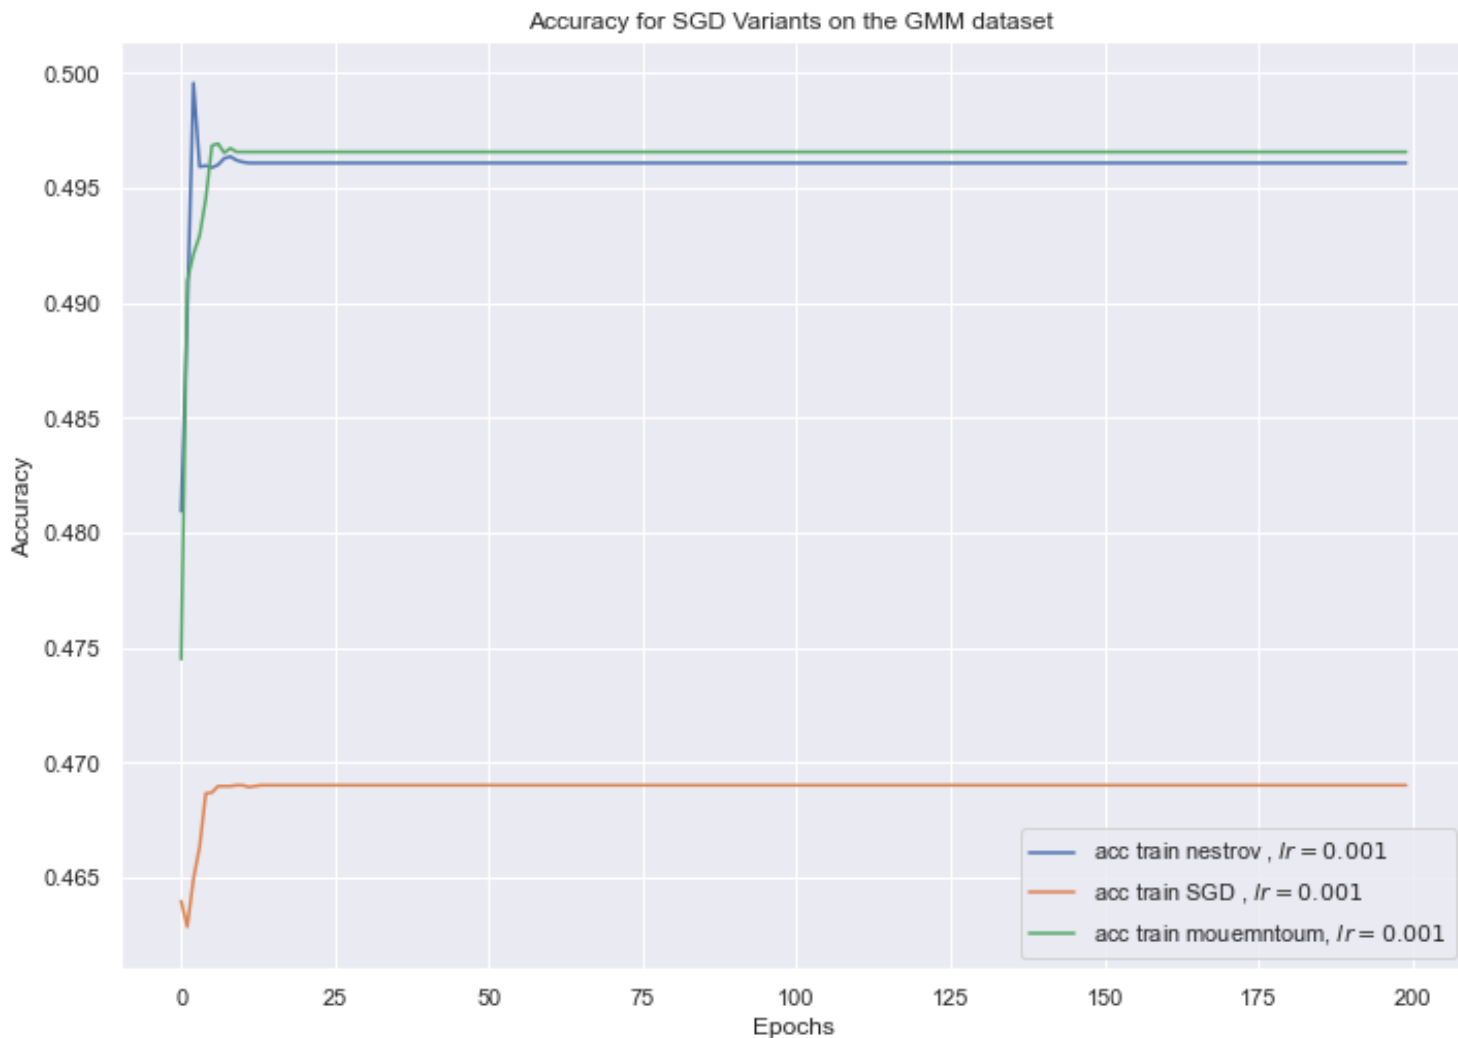
2 y_train=mat["Ct"]
4 x_train=mat["Yt"]

6 y_val=mat["Cv"]
  x_val=mat["Yv"]

8
10 x_train=np.vstack([x_train,np.ones(25000)])
    x_val=np.vstack([x_val,np.ones(6250)])

```

```
comapre_sgd_variants(data_set_peaks=False)
```



Learning rate and batch sizes experiments for peaks dataset

Varying batch sizes still resulted an almost equal accuracy in convergence if we adjust the lr accordingly, but if the lr is constant and the batch sizes are changing then there are differences. for example, if $lr=0.001$

for batchsize 16 we reach convergence with 53.6% for batchsize 32 we reach convergence with 53.6% for batchsize 128 we reach convergence with 51% for batchsize 256 we reach convergence with 50%

but for $lr=0.01$ we converge with batch-size 256 to 53.6%, which make sense since we're averaging 256 samples which should grant us a good gradient, therefore we can use higher learning rate since we're more confident.

eventually, we chose a batchsize of 32 and examined the following learning rates :

$lr=0.1$, convergence to 58.6 and less stable (upper limit=0.59) $lr=0.01$ convergence to 58.6, stability (upper limit =0.545) $lr=0.001$, convergence to 58.6 and stability (upper limit =0.586) $lr=0.0001$, convergence to 49.7 stability (upper limit =0.4975)

therefor learning rate impacts how the accuracy fluctuates around its around convergence , which could result in unstable learning. Learning rates that are too low would hurt our performance. As learning rate decreases the stability of the training necessary increases but the accuracy changes and stagants if the lr is too low.

GMM dataset

As we can see , the accuracy is lower at convergence , the added dimensions of the input heavily complicates the space we wish to learn , the peaks data set is a function from $f : \mathbb{R}^2 \rightarrow \mathbb{R}^5$, but the GMM datasets increase the dimension of the input to 5 , i.e. $f : \mathbb{R}^5 \rightarrow \mathbb{R}^5$ which explains the different in accuracy.

Part 2

The following section defines the forward and backwards of a standart feed forward network using tanh activation function.

```

import numpy as np
2
from copy import deepcopy
4
def ReLU(x):
6     return x * (x > 0)

8 def softmax(z):
    z -= np.max(z)
10    sm = (np.exp(z).T / np.sum(np.exp(z),axis=1)).T
    return sm
12
def forward_FFN(x, y, W, b, act_f):
14    depth=len(W)
    hidden_layers=[x]
16    hidden_layer=x
    for i in range(depth):
18        if i!=depth-1:
            hidden_layer=act_f(hidden_layer@W[i]+b[i])
20            hidden_layers.append(hidden_layer)
        else:
22            #no act_f on the final layer
            hidden_layer =(hidden_layer @ W[i] + b[i])
24            hidden_layers.append(hidden_layer)

26
    exp_scores = np.exp(deepcopy(hidden_layer))
28    probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)
    loss = (-1/x.shape[0]) * np.sum(y * np.log(probs))
30    return probs , hidden_layers ,loss
32
34
def backward_FFN(x, y, W, b):
36    num_examples=x.shape[0]
    prob,hidden_layers,loss=forward_FFN(x, y, W, b, np.tanh)
38
    #derv of softmax
40    dscores = deepcopy(prob)
    dscores[range(num_examples),np.argmax(y,axis=1)]-=1
42    dscores /= num_examples
    depth=len(W)
44

```

```

46     V=dscores
48
49     dW_list=[]
50     dB_list=[]
51     dX_list=[V]
52     curr_V=V
53     for j in range(depth-1,-1,-1):
54         curr_w=W[j]
55         curr_layer=hidden_layers[j]
56         dW=curr_layer.T@curr_V
57         db=np.sum(curr_V,axis=0)
58         dXn=curr_V@curr_w.T
59         V=dXn
60
61
62         if j!=0:
63             dtanh = 1 - np.tanh(hidden_layers[j-1]@W[j-1]+b[j-1]) ** 2
64             curr_V = dtanh * V
65
66
67         dW_list.append(dW)
68         dB_list.append(db)
69         dX_list.append(curr_V)
70
71     return dW_list[::-1],dB_list[::-1],dX_list[::-1]
72
73 import scipy.io
74 mat = scipy.io.loadmat('PeaksData.mat')

```

Jacobain test feed foward network

We perform the jacobian test on a nernal network with 2 layer , the structure of the testes network is as follows:

$$\text{softmax}(W_2 \tanh(W_1 x_0 + b_1) + b_2)$$

thus , if the depth increases then the recurring part is the tanh layer , thus we perform a jacobian test on that layer. The last layer , the softmax and the loss function would be verified afterwards using the gradient test on the whole feed forward network.

Part 2 Question 1

Jacboain test for feed forward

```

def Jacobian_test_feed_forward_NN():
2     """
3     The first part of the function defines random weights and biases , then we perform the
4     jacobain test
5     for the first layer in the network.
6     :return:
7     """
8     values_linear=[]
9     values_quard=[]
10    epsilon=0.2
11    #wieghts init
12    w_orig=[np.random.random((5,20)),np.random.random((20,5))]
13    b_orig=[np.random.random(20),np.random.random(5)]
14
15    x=np.random.random((1,5))
16    y=np.zeros((1,5))

```



```

16     y[:,3]=1
18     prob, layers, loss = forward_FFN(x, y, w_orig, b_orig, np.tanh)
19     #random vectors init
20     random_w1, random_w2 = np.random.random((5, 20)), np.random.random((20, 5))
21     random_w1 = random_w1 / np.linalg.norm(random_w1, ord=2)
22
23     random_b1, random_b2 = np.random.random(20), np.random.random(5)
24     w = deepcopy(w_orig)
25     b = deepcopy(b_orig)
26
27     for i in range(20):
28         eps = epsilon * (np.power(0.5, i))
29         w[0] = w_orig[0] + random_w1 * eps
30         b[0] = b_orig[0] + random_b1 * eps
31         prob, layers_eps, loss = forward_FFN(x, y, w, b, np.tanh)
32
33         #jacobain calculation for the first layer.
34         derv_first_layer = 1 - np.tanh((x @ w_orig[0]) + b_orig[0]) ** 2
35         grad_w = x.T @ derv_first_layer
36         grad_b = (derv_first_layer * eps * random_b1).sum(axis=0)
37         grad_mult_random_w = ((grad_w * random_w1 * eps).sum(axis=0, keepdims=True))
38
39         #quad values calcuation
40         grad_test_quard = layers_eps[1] - layers[1] - grad_mult_random_w - grad_b
41         val_linear = np.linalg.norm(layers_eps[1] - layers[1], ord=2)
42         val_quad = np.linalg.norm(grad_test_quard, ord=2)
43         values_linear.append(val_linear)
44         values_quard.append(val_quad)
45
46     values_quard = np.array(values_quard)
47     values_linear = np.array(values_linear)
48
49     return values_linear, values_quard
50
51 import seaborn as sns

```


The following section prints the results in two lists, the first column signifies the quadratic or linear measurements and the second column displays the ration between following outputs, as we can see the ratio for the quadratic converges to 4, while the linear convergences to 2 which proves that our Jacobian is correct.

```

2 values_linear, values_quard = Jacobian_test_feed_forward_NN()
3 quad_graident_test = pd.DataFrame(data=[values_quard, values_quard[:-1] / values_quard[1:]]) .T
4 quad_graident_test.columns = ["quadratic", "F(i)/F(i+1)"]
5 linear_gradient_test = pd.DataFrame(data=[values_linear, values_linear[:-1] / values_linear[1:]]) .T
6 linear_gradient_test.columns = ["linear", "F(i)/F(i+1)"]
7
8 print(quad_graident_test)
9 print(linear_gradient_test)
10
11 plt.plot(values_quard)
12 plt.plot(values_linear)
13 plt.yscale("log")
14 plt.title("Jacobin test for feed forward $tanh(W_{0}x+b_{0})$")
15 plt.ylabel("Gradient test values")
16 plt.xlabel("Iterations")
17 plt.show()
18
19 quadratic    F(i)/F(i+1)

```

0	1.839252e-02	3.794304
3 1	4.847404e-03	3.894720
2	1.244609e-03	3.946777
5 3	3.153482e-04	3.973246
4	7.936791e-05	3.986588
7 5	1.990873e-05	3.993285
6	4.985551e-06	3.996641
9 7	1.247436e-06	3.998320
8	3.119899e-07	3.999160
11 9	7.801387e-08	3.999580
10	1.950552e-08	3.999790
13 11	4.876635e-09	3.999895
12	1.219191e-09	3.999948
15 13	3.048016e-10	3.999971
14	7.620095e-11	3.999994
17 15	1.905027e-11	4.000030
16	4.762532e-12	3.999889
19 17	1.190666e-12	4.000523
18	2.976276e-13	4.002611
21 19	7.435837e-14	NaN
linear F(i)/F(i+1)		
23 0	1.012884e-01	1.844962
1	5.490000e-02	1.919358
25 2	2.860331e-02	1.958883
3	1.460185e-02	1.979241
27 4	7.377498e-03	1.989570
5	3.708086e-03	1.994773
29 6	1.858902e-03	1.997383
7	9.306686e-04	1.998691
31 8	4.656391e-04	1.999345
9	2.328958e-04	1.999673
33 10	1.164670e-04	1.999836
11	5.823825e-05	1.999918
35 12	2.912032e-05	1.999959
13	1.456046e-05	1.999980
37 14	7.280303e-06	1.999990
15	3.640170e-06	1.999995
39 16	1.820090e-06	1.999997
17	9.100461e-07	1.999999
41 18	4.550233e-07	1.999999
19	2.275117e-07	NaN



Part 2 Question 2

Forward and backwards ResNet

```
def RNN_forward(x,y,W,b,act_f):
2   depth=len(W)
   hidden_layers=[x]
4   hidden_layer=x.T
   hidden_tanh=[]
6   hidden_relu=[]
   for i in range(depth):
8       #Forward caculation of ReLU(w1x+tanh(w2x+b1)+b2)
       hidden_layer_1 = W[i][0] @ hidden_layer + b[i][0].reshape(-1,1)
10      hidden_tanh.append(W[i][0] @ hidden_layer + b[i][0].reshape(-1,1))
       hidden_relu.append((W[i][1]@ hidden_layer + (act_f(hidden_layer_1)) +
12      b[i][1].reshape(-1,1)))
       hidden_layer= ReLU((W[i][1]@ hidden_layer + (act_f(hidden_layer_1)) +
14      b[i][1].reshape(-1,1)))
       hidden_layers.append(hidden_layer.T)

   exp_scores = np.exp(deepcopy(((hidden_layer.T))))
16  probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)
   loss = (-1/x.shape[0]) * np.sum(y* np.log(probs))
18  return probs , hidden_layers ,loss,hidden_relu,hidden_tanh
```

```

20 def backward_rnn(x,y,W,b):
21     num_examples=x.shape[0]
22     prob,hidden_layers,loss,hidden_relu,hidden_tanh=RNN_forward(x,y,W,b,np.tanh)
23
24     dscores = deepcopy(prob)
25     dscores[range(num_examples),np.argmax(y,axis=1)]-=1
26     dscores /= num_examples
27     depth=len(W)
28     V=dscores
29
30
31     dW_list=[]
32     dB_list=[]
33     dX_list=[]
34
35     for j in range(depth-1,-1,-1):
36         w1=W[j][0]
37         w2=W[j][1]
38         b1=b[j][0]
39         b2=b[j][1]
40
41         drelu=hidden_relu[j].T
42
43         drelu[drelu<0]=0
44         drelu[drelu>0]=1
45
46         V = drelu * V
47         X=hidden_layers[j]
48         w_outer=(X.T@V).T
49         db_outer=np.sum(V,axis=0)
50         dhidden=w1 @ X.T + b1.reshape(-1,1)
51         dhidden=1-np.tanh(dhidden)*np.tanh(dhidden)
52
53         dXn= V@w2 + (dhidden.T*V)@w1
54
55         w_inner= ((V.T*dhidden)@ X)
56
57         db_inner= (dhidden *V.T).sum(axis=1)
58
59         V=dXn
60         dW_list.append([w_inner,w_outer])
61         dB_list.append([db_inner,db_outer])
62         dX_list.append(dXn)
63
64     return dW_list[::-1],dB_list[::-1],dX_list[::-1]

```

Jacobin test for ResNet

Results as shown as a table , one for the quadratic and one for the linear as in previous examples.

```

2
3 def init_weights_forward_ResNet(depth,WIDTH_NET,ORDER_NORM,NUM_CLASSES,INPUT_LEN):
4
5     w=[[np.random.random((WIDTH_NET, INPUT_LEN))/ORDER_NORM,np.random.random((WIDTH_NET,
6     INPUT_LEN))/ORDER_NORM]]
7     for i in range(depth-2):
8         w.append([np.random.random((WIDTH_NET, WIDTH_NET)) / ORDER_NORM,
9         np.random.random((WIDTH_NET, WIDTH_NET))/ORDER_NORM])

```

```

8     w.append([np.random.random((NUM_CLASSES,WIDTH_NET))/ORDER_NORM,np.random.random((NUM_CLASSES,W
10
11     b=[np.random.random(WIDTH_NET)/ORDER_NORM,np.random.random(WIDTH_NET)/ORDER_NORM]]
12     for i in range(depth-2):
13         b.append([np.random.random(WIDTH_NET)/ORDER_NORM,np.random.random(WIDTH_NET)/ORDER_NORM])
14     b.append([np.random.random(NUM_CLASSES) / ORDER_NORM,np.random.random(NUM_CLASSES) /
15         ORDER_NORM])
16     return w,b
17
18 def Jacobian_test_ResNet():
19     """
20     The first part of the function defines random weights and biases , then we perform the
21         jacobain test
22     for the first layer in the network.
23     :return:
24     """
25     values_linear=[]
26     values_quard=[]
27     epsilon=0.2
28
29     #initilziaotn of weights for the resnet network
30     w_orig,b_orig=init_weights_forward_ResNet(2,WIDTH_NET=20,ORDER_NORM=1,NUM_CLASSES=5,INPUT_LEN=
31
32     x=np.random.random((1,5))
33     y=np.zeros((1,5))
34     y[:,3]=0
35
36     # init of random vectors
37     random_w00,random_w01=np.random.random((20,5)),np.random.random((20,5))
38     random_b1,random_b2=np.random.random(20),np.random.random(20)
39     w=deepcopy(w_orig)
40     b=deepcopy(b_orig)
41
42     #hand calculation of the resnet layer
43     x=x.T
44     starting_layer=(w[0][0]@x+b[0][0].reshape(-1,1))
45     middle_layer=w[0][1]@x+np.tanh(w[0][0]@x+b[0][0].reshape(-1,1))+b[0][1].reshape(-1,1)
46     final_layer=np.tanh(w[0][1]@x+np.tanh(w[0][0]@x+b[0][0].reshape(-1,1))+b[0][1].reshape(-1,1))
47
48     #derv of tanh
49     outer_derv=1-np.tanh(middle_layer)**2
50
51     #derv of w1
52     w01_grad=x@outer_derv.T
53     w01_grad=w01_grad.T
54
55
56     inner_derv=1-np.tanh(starting_layer)**2
57
58     V=inner_derv*outer_derv
59
60     #derv of w0
61     w00_derv=x@V.T
62     w00_derv=w00_derv.T
63
64

```

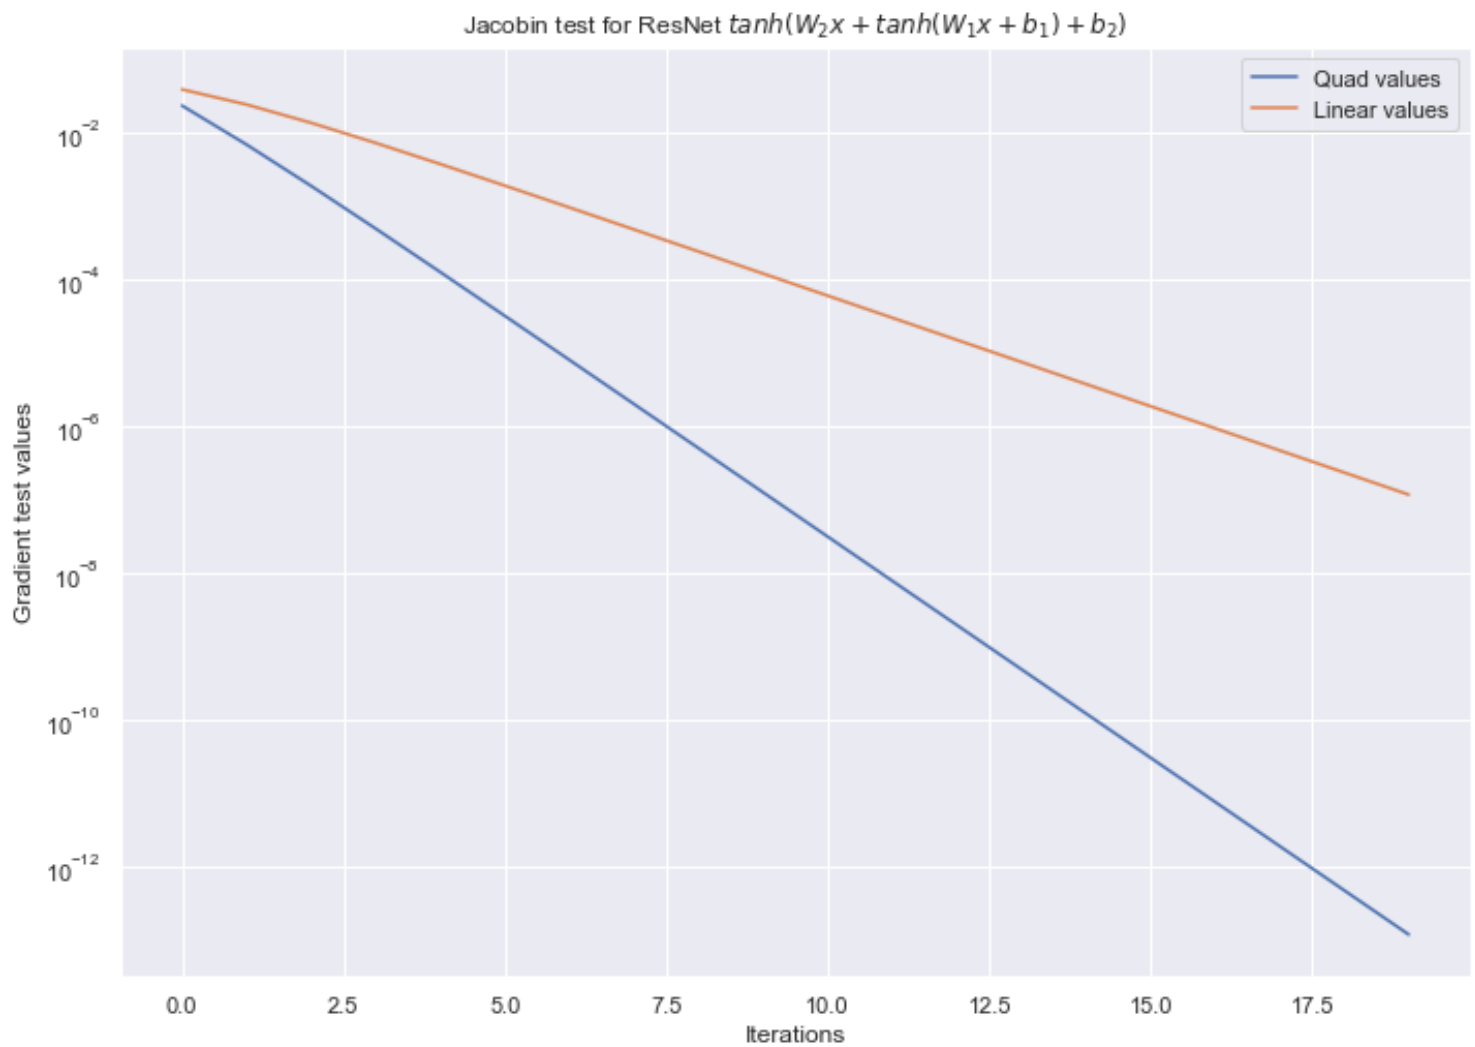
```

66     # calcuation of the biases gradients
67     b2_derv=outer_derv.sum(axis=1,keepdims=True)
68     b1_derv=V.sum(axis=1,keepdims=True)
69
70     b2_derv=b2_derv.T
71     b1_derv=b1_derv.T
72
73     for i in range(20):
74         eps = epsilon * (np.power(0.5, i))
75         w[0][1]=w_orig[0][1]+random_w01*eps
76         w[0][0]=w_orig[0][0]+random_w00*eps
77         b[0][1]=b_orig[0][1]+random_b2*eps
78         b[0][0]=b_orig[0][0]+random_b1*eps
79         #jacobain calculation for the first layer.
80
81
82         final_layer_moved=np.tanh(w[0][1]*x+np.tanh(w[0][0]*x+b[0][0].reshape(-1,1))+b[0][1].resha
83         grad_w1 = (w01_grad * random_w01 * eps).sum(axis=1, keepdims=True)
84         grad_w0 = (w00_derv * random_w00 * eps).sum(axis=1, keepdims=True)
85
86         grad_b2 = b2_derv * random_b2 * eps
87         grad_b1=b1_derv*random_b1*eps
88
89
90         grad_test_quard = final_layer_moved - final_layer - grad_w1 - grad_w0 -
91             grad_b2.T-grad_b1.T
92         val_linear = np.linalg.norm(final_layer_moved-final_layer,ord=2)
93         val_quad = np.linalg.norm(grad_test_quard, ord=2)
94         values_linear.append(val_linear)
95         values_quard.append(val_quad)
96
97         values_quard=np.array(values_quard)
98         values_linear=np.array(values_linear)
99
100     return values_linear,values_quard
101
102 values_linear,values_quard=Jacobian_test_ResNet()
103 quad_graident_test=pd.DataFrame(data=[values_quard,values_quard[:-1]/values_quard[1:]]).T
104 quad_graident_test.columns=["quadratic","F(i)/F(i+1)"]
105 linear_gradient_test=pd.DataFrame(data=[values_linear,values_linear[:-1]/values_linear[1:]]).T
106 linear_gradient_test.columns=["linear","F(i)/F(i+1)"]
107 print(quad_graident_test)
108 print(linear_gradient_test)
109 plt.plot(values_quard,label="Quad values")
110 plt.plot(values_linear,label="Linear values")
111
112 plt.legend()
113 plt.yscale("log")
114 plt.title("Jacobin test for ResNet $tanh(W_2x+tanh(W_1x+b_1)+b_2)$")
115 plt.ylabel("Gradient test values")
116 plt.xlabel("Iterations")
117 plt.show()

```

	quadratic	F(i)/F(i+1)
0	2.291354e-02	3.394476
3 1	6.750244e-03	3.654941
2	1.846882e-03	3.814575

5	3	4.841644e-04	3.903693
	4	1.240273e-04	3.950895
7	5	3.139221e-05	3.975202
	6	7.897008e-06	3.987539
9	7	1.980422e-06	3.993754
	8	4.958797e-07	3.996873
11	9	1.240669e-07	3.998436
	10	3.102887e-08	3.999217
13	11	7.758735e-09	3.999609
	12	1.939873e-09	3.999804
15	13	4.849921e-10	3.999902
	14	1.212510e-10	3.999950
17	15	3.031313e-11	3.999999
	16	7.578284e-12	3.999954
19	17	1.894593e-12	3.999963
	18	4.736526e-13	4.002315
21	19	1.183447e-13	NaN
		linear	F(i)/F(i+1)
23	0	3.806945e-02	1.607284
	1	2.368558e-02	1.774717
25	2	1.334611e-02	1.878673
	3	7.104010e-03	1.936937
27	4	3.667651e-03	1.967836
	5	1.863799e-03	1.983756
29	6	9.395304e-04	1.991837
	7	4.716905e-04	1.995908
31	8	2.363288e-04	1.997951
	9	1.182855e-04	1.998975
33	10	5.917309e-05	1.999487
	11	2.959413e-05	1.999744
35	12	1.479896e-05	1.999872
	13	7.399956e-06	1.999936
37	14	3.700096e-06	1.999968
	15	1.850078e-06	1.999984
39	16	9.250463e-07	1.999992
	17	4.625250e-07	1.999996
41	18	2.312630e-07	1.999998
	19	1.156316e-07	NaN



```

import scipy.io
2 mat = scipy.io.loadmat('PeaksData.mat')

4 y=mat["Ct"].T
  x=mat["Yt"].T
6 x=x[:1,:]
  y=y[:1,:]
8
def grad_test_whole_network():
10     values_linear=[]
    values_quard=[]
12
    NUM_CLASSES = 5
14     INPUT_LEN=2
    WIDTH_NET = 50
16
    # Init random weights
18     w1= np.random.random((INPUT_LEN, WIDTH_NET))

20     w4= np.random.random((WIDTH_NET, NUM_CLASSES))
    b1= np.random.random(WIDTH_NET)
22
    b4= np.random.random(NUM_CLASSES)
24
    w=[w1,w4]

```



```

26     b=[b1,b4]
    #init random vectors
28     w1_random=np.random.random((INPUT_LEN,WIDTH_NET ))
    w1_random= w1_random/np.linalg.norm(w1_random,ord=2)
30
    w2_random=np.random.random((WIDTH_NET,NUM_CLASSES))
32     w2_random= w2_random/np.linalg.norm(w2_random,ord=2)
34
    b1_random=np.random.random(WIDTH_NET)
    b1_random= b1_random/np.linalg.norm(b1_random,ord=2)
36
    b2_random=np.random.random(NUM_CLASSES)
38     b2_random=b2_random/np.linalg.norm(b2_random,ord=2)
40
    probs,_,loss=forward_FFN(x,y,w,b,np.tanh)
42
    dw,db,dX=backward_FFN(x,y,w,b)
    eps=0.3
44     w_test=deepcopy(w)
    b_test=deepcopy(b)
46     for i in range(10):
        #add random vectors to weights
48         epsilom = eps * (np.power(0.5, i))
        w_test[0]=w[0]+w1_random*epsilom
50         w_test[1]=w[1]+w2_random*epsilom
        b_test[0]=b[0]+b1_random*epsilom
52         b_test[1]=b[1]+b2_random*epsilom
54
        # grad summation of the final values for the grad test in the quad values
        grad=(dw[0]*w1_random*epsilom).sum()+(dw[1]*w2_random*epsilom).sum()+(db[0]*b1_random*epsilom).sum()
56
        probs,_,loss_moved=forward_FFN(x,y,w_test,b_test,np.tanh)
58         val_linear = np.abs(loss_moved-loss)
        val_quad = np.abs(loss_moved- loss - grad)
60
        values_linear.append(val_linear)
62         values_quard.append(val_quad)
64
        values_quard=np.array(values_quard)
66         values_linear=np.array(values_linear)
        return values_quard,values_linear
68
    values_quard,values_linear=grad_test_whole_network()
70     quad_graident_test=pd.DataFrame(data=[values_quard,values_quard[:-1]/values_quard[1:]]).T
    quad_graident_test.columns=["quadratic","F(i)/F(i+1)"]
72     linear_gradient_test=pd.DataFrame(data=[values_linear,values_linear[:-1]/values_linear[1:]]).T
    linear_gradient_test.columns=["linear","F(i)/F(i+1)"]
74     display(quad_graident_test)
    display(linear_gradient_test)

```

```

1         quadratic  F(i)/F(i+1)
0  1.479319e-03      3.817278
3  1  3.875326e-04      3.914214
2  9.900648e-05      3.958446
5  3  2.501145e-05      3.979551
4  6.284992e-06      3.989857
7  5  1.575242e-06      3.994949
6  3.943086e-07      3.997479

```

```

9 7 9.863930e-08 3.998741
8 2.466759e-08 3.999371
11 9 6.167868e-09 NaN

```

```

1 linear F(i)/F(i+1)
0 0.165942 2.008524
3 1 0.082619 2.004598
2 0.041215 2.002380
5 3 0.020583 2.001210
4 0.010285 2.000610
7 5 0.005141 2.000306
6 0.002570 2.000153
9 7 0.001285 2.000077
8 0.000642 2.000038
11 9 0.000321 NaN

```

Part 2 , Question 3 , Gradient test feed forward

```

import scipy.io
2 mat = scipy.io.loadmat('PeaksData.mat')

4 y=mat["Ct"].T
  x=mat["Yt"].T
6 x=x[:1,:]
  y=y[:1,:]
8 import numpy as np
def grad_test_whole_network():
10     values_linear=[]
    values_quard=[]

12
    NUM_CLASSES = 5
14     INPUT_LEN=2
    WIDTH_NET = 50

16
    #init weights for the netowkr

18
    w1= np.random.random((INPUT_LEN, WIDTH_NET))

20
    w4= np.random.random((WIDTH_NET, NUM_CLASSES))
22     b1= np.random.random(WIDTH_NET)

24
    b4= np.random.random(NUM_CLASSES)

26
    w=[w1,w4]
    b=[b1,b4]
28     #init random vectors
    w1_random=np.random.random((INPUT_LEN,WIDTH_NET ))
30     w1_random= w1_random/np.linalg.norm(w1_random,ord=2)

32
    w2_random=np.random.random((WIDTH_NET,NUM_CLASSES))
    w2_random= w2_random/np.linalg.norm(w2_random,ord=2)

34
    b1_random=np.random.random(WIDTH_NET)
36     b1_random= b1_random/np.linalg.norm(b1_random,ord=2)

38
    b2_random=np.random.random(NUM_CLASSES)
    b2_random=b2_random/np.linalg.norm(b2_random,ord=2)
40

```

```

42     probs,_,loss=forward_FFN(x,y,w,b,np.tanh)

44     dw,db,dX=backward_FFN(x,y,w,b)
     eps=0.3
46     w_test=deepcopy(w)
     b_test=deepcopy(b)
48     for i in range(10):
         #add random vectors to the weights t
50         epsilom = eps * (np.power(0.5, i))
         w_test[0]=w[0]+w1_random*epsilom
52         w_test[1]=w[1]+w2_random*epsilom
         b_test[0]=b[0]+b1_random*epsilom
54         b_test[1]=b[1]+b2_random*epsilom

56         #final values of the gradient for the grad test
         grad=(dw[0]*w1_random*epsilom).sum()+(dw[1]*w2_random*epsilom).sum()+(db[0]*b1_random*epsilom).sum()

58         probs,_,loss_moved=forward_FFN(x,y,w_test,b_test,np.tanh)
         val_linear = np.abs(loss_moved-loss)
         val_quad = np.abs(loss_moved- loss - grad)

62         values_linear.append(val_linear)
64         values_quard.append(val_quad)

66
68     values_quard=np.array(values_quard)
     values_linear=np.array(values_linear)
     return values_quard,values_linear

70
values_quard,values_linear=grad_test_whole_network()
72 quad_gradient_test=pd.DataFrame(data=[values_quard,values_quard[:-1]/values_quard[1:]]).T
quad_gradient_test.columns=["quadratic","F(i)/F(i+1)"]
74 linear_gradient_test=pd.DataFrame(data=[values_linear,values_linear[:-1]/values_linear[1:]]).T
linear_gradient_test.columns=["linear","F(i)/F(i+1)"]
76 display(quad_gradient_test)
display(linear_gradient_test)

78
plt.plot(values_quard,label="Quad values")
80 plt.plot(values_linear,label="Linear values")

82 plt.legend()
plt.yscale("log")
84 plt.title("Gradient test for feed forward network")
plt.ylabel("Gradient test values")
86 plt.xlabel("Iterations")
plt.show()


```

	quadratic	F(i)/F(i+1)
0	8.002958e-03	4.055892
3 1	1.973168e-03	4.028968
2	4.897454e-04	4.014711
5 3	1.219877e-04	4.007408
4	3.044055e-05	4.003717
7 5	7.603072e-06	4.001862
6	1.899884e-06	4.000932
9 7	4.748604e-07	4.000466
8	1.187013e-07	4.000233
11 9	2.967359e-08	NaN

```

1      linear  F(i)/F(i+1)
0  0.096207    2.088044
3  1  0.046075    2.044084
2  0.022541    2.022047
5  3  0.011147    2.011024
4  0.005543    2.005512
7  5  0.002764    2.002756
6  0.001380    2.001378
9  7  0.000690    2.000689
8  0.000345    2.000344
11 9  0.000172    NaN

```



Part 2 Question 3 , Gradient test ResNet

```

1 mat = scipy.io.loadmat('PeaksData.mat')
  y=mat["Ct"].T
3 x=mat["Yt"].T
  x=x[:1,:]
5 y=y[:1,:]
def grad_test_whole_network():
7     values_linear=[]
  values_quard=[]
9
11     NUM_CLASSES = 5

```

```

13 INPUT_LEN=2
14 WIDTH_NET = 100
15 SCALE=1
16
17 #init of random weights for the network
18 w00 = np.random.random((WIDTH_NET, INPUT_LEN))/SCALE
19 w01 = np.random.random((WIDTH_NET, INPUT_LEN))/SCALE
20 w10 = np.random.random(( NUM_CLASSES,WIDTH_NET))/SCALE
21 w11 = np.random.random((NUM_CLASSES,WIDTH_NET))/SCALE
22 b00 = np.random.random(WIDTH_NET)/SCALE
23 b01 = np.random.random(WIDTH_NET)/SCALE
24 b10 = np.random.random(NUM_CLASSES)/SCALE
25 b11 = np.random.random(NUM_CLASSES)/SCALE
26
27 w = [[w00, w01],[w10,w11]]
28 b = [[b00, b01],[b10,b11]]
29
30 #init of random vectors to add to the weights
31 w00_random=np.random.random((WIDTH_NET,INPUT_LEN ))
32 w00_random= w00_random/np.linalg.norm(w00_random,ord=2)
33
34 w01_random=np.random.random((WIDTH_NET,INPUT_LEN))
35 w01_random= w01_random/np.linalg.norm(w01_random,ord=2)
36
37 w10_random=np.random.random((NUM_CLASSES,WIDTH_NET))
38 w10_random= w10_random/np.linalg.norm(w10_random,ord=2)
39
40 w11_random=np.random.random((NUM_CLASSES,WIDTH_NET))
41 w11_random= w11_random/np.linalg.norm(w11_random,ord=2)
42
43 b00_random=np.random.random(WIDTH_NET)
44 b00_random= b00_random/np.linalg.norm(b00_random,ord=2)
45
46 b01_random=np.random.random(WIDTH_NET)
47 b01_random= b01_random/np.linalg.norm(b01_random,ord=2)
48
49 b10_random=np.random.random(NUM_CLASSES)
50 b10_random= b10_random/np.linalg.norm(b10_random,ord=2)
51
52 b11_random=np.random.random(NUM_CLASSES)
53 b11_random= b11_random/np.linalg.norm(b11_random,ord=2)
54
55
56 probs,_,loss,_,_=RNN_forward(x,y,w,b,np.tanh)
57
58 dw,db,dX=backward_rnn(x,y,w,b)
59 eps=0.3
60 w_test=deepcopy(w)
61 b_test=deepcopy(b)
62 for i in range(10):
63     epsilon = eps * (np.power(0.5, i))
64
65     #adding the random vectors to the weights accordinly
66     w_test[1][1]=w[1][1]+w11_random*epsilon
67     w_test[1][0]=w[1][0]+w10_random*epsilon
68     w_test[0][1]=w[0][1]+w01_random*epsilon
69     w_test[0][0]=w[0][0]+w00_random*epsilon

```

```

71     b_test[1][1]=b[1][1]+b11_random*epsilon
       b_test[1][0]=b[1][0]+b10_random*epsilon
73     b_test[0][1]=b[0][1]+b01_random*epsilon
       b_test[0][0]=b[0][0]+b00_random*epsilon
75
77     #calculation of the gradient per layer n their summation with regard to b and w
       grad_w=(dw[1][1]*w11_random*epsilon).sum()+(dw[0][0]*w00_random*epsilon).sum()+(dw[0][1]*w
79     grad_b=(db[1][1]*b11_random*epsilon).sum()+(db[0][0]*b00_random*epsilon).sum()+(db[0][1]*b
       grad=grad_b+grad_w
81
       probs,_,loss_moved,_,_=RNN_forward(x,y,w_test,b_test,np.tanh)
83     val_linear = np.abs(loss_moved-loss)
       val_quad = np.abs(loss_moved- loss - grad)
85
       values_linear.append(val_linear)
87     values_quard.append(val_quad)
89
       values_quard=np.array(values_quard)
91     values_linear=np.array(values_linear)
       return values_quard,values_linear
93
values_quard,values_linear=grad_test_whole_network()
95 quad_graident_test=pd.DataFrame(data=[values_quard,values_quard[:-1]/values_quard[1:]]).T
       quad_graident_test.columns=["quadratic","F(i)/F(i+1)"]
97 linear_gradient_test=pd.DataFrame(data=[values_linear,values_linear[:-1]/values_linear[1:]]).T
       linear_gradient_test.columns=["linear","F(i)/F(i+1)"]
99 display(quad_graident_test)
       display(linear_gradient_test)
101
       plt.plot(values_quard,label="Quad values")
103 plt.plot(values_linear,label="Linear values")
       plt.legend()
105 plt.yscale("log")
       plt.title("Gradient test for ResNet full network")
107 plt.ylabel("Gradient test values")
       plt.xlabel("Iterations")
109 plt.show()

```

```

1      quadratic  F(i)/F(i+1)
0  1.384215e-02    9.490286
3  1  1.458560e-03    3.294475
2  4.427289e-04    4.004528
5  3  1.105571e-04    4.002310
4  2.762331e-05    4.001167
7  5  6.903815e-06    4.000586
6  1.725701e-06    4.000294
9  7  4.313935e-07    4.000147
8  1.078444e-07    4.000073
11 9  2.696061e-08           NaN


```

```

1      linear  F(i)/F(i+1)
0  0.122617    1.836380
3  1  0.066771    1.982980
2  0.033672    1.986923
5  3  0.016947    1.993494
4  0.008501    1.996755
7  5  0.004257    1.998379

```

6	0.002130	1.999190	
9	7	0.001066	1.999595
8	0.000533	1.999798	
11	9	0.000266	NaN



Part 2 Question 4 , optimization of feed forward neural network on Peaks dataset

After verifying we perform a simple optimization of the neural network using the given datasets , peaks and GMM.

```

1
3 import matplotlib.pyplot as plt
5 def init_weights_forward_NN(depth,WIDTH_NET,ORDER_NORM,NUM_CLASSES,INPUT_LEN):
7     # init w matrices
8     w1= np.random.random((INPUT_LEN, WIDTH_NET)) / ORDER_NORM
9     w=[w1]
10    for i in range(depth-2):
11        w.append(np.random.random((WIDTH_NET, WIDTH_NET)) / ORDER_NORM)
12        w.append(np.random.random((WIDTH_NET, NUM_CLASSES)) / ORDER_NORM)
13
14    print(f"the amount of layers is {len(w)}")
15
16    #init biases

```

```

17     b1= np.random.random(WIDTH_NET) / ORDER_NORM
18     b=[b1]
19     for i in range(depth-2):
20         b.append(np.random.random(WIDTH_NET) / ORDER_NORM)
21     b.append(np.random.random(NUM_CLASSES) / ORDER_NORM)
22     return w,b
23
24
25
26
27
28
29 def
optimize_NN(x,y,x_val,y_val,depth=4,lr=0.01,plot=True,batch_size=16,data_set="Peaks",iters=150)
    WIDTH_NET = net_width
30     # init network sizes based on dataset
31     if data_set=="Peaks":
32         ORDER_NORM = 1
33         NUM_CLASSES = 5
34         INPUT_LEN=2
35     if data_set=="GMM":
36         ORDER_NORM = 1
37         NUM_CLASSES = 5
38         INPUT_LEN=5
39
40
41
42
43     w,b=init_weights_forward_NN(depth=depth,WIDTH_NET=WIDTH_NET,ORDER_NORM=ORDER_NORM,NUM_CLASSES=
44
45
46     loss_train=[]
47     loss_val=[]
48     acc_train=[]
49     acc_val=[]
50     train_size = 25000
51     val_size=6250
52     #each iter is an epoch
53     for i in range(iters):
54         curr_lr=np.power(0.99,i)*lr
55         #each loop is per batch with weights update
56         for j in range(train_size // batch_size):
57             indices=np.random.randint(0, train_size, batch_size)
58             dw,db,dx=backward_FFN(x[indices, :], y[indices, :], w, b)
59
60             for i in range(len(w)):
61                 b[i]=b[i]-curr_lr*db[i]
62                 w[i]=w[i]-curr_lr*dw[i]
63
64
65         #calc accuracy on train
66         probs, hl, loss = forward_FFN(x[:, :], y[:, :], w, b, np.tanh)
67         agreements = np.argmax(probs, axis=1) == np.argmax(y[:, :], axis=1)
68         acc_train.append(agreements.sum() / train_size)
69         loss_train.append(loss)
70
71         #calc accuracy on val
72         probs, hl, loss = forward_FFN(x_val[:, :], y_val[:, :], w, b, np.tanh)
73         agreements = np.argmax(probs, axis=1) == np.argmax(y_val[:, :], axis=1)
74         acc_val.append(agreements.sum() / val_size)

```



```

75     loss_val.append(loss)
77
79
80     if plot:
81         plt.plot(acc_train, label="acc train")
82         plt.plot(acc_val, label="acc val")
83         plt.legend()
84         plt.title(f"Accuracy on {data_set} data set , using SGD")
85         plt.xlabel("Epochs")
86         plt.ylabel("Accuracy")
87         print(f"Final accuracy on train is {acc_train[-1]}")
88         print(f"Final accuracy on val is {acc_val[-1]}")
89         plt.show()
90
91         plt.plot(loss_train, label="loss train")
92         plt.plot(loss_val, label="loss val")
93         plt.legend()
94         plt.title("Loss on Peaks data set , using SGD")
95         plt.yscale("log")
96         plt.xlabel("Epochs")
97         plt.ylabel("Loss")
98         plt.show()
99
100     plt.show()
101     return w, loss , acc_val, acc_train

```

```


1 import scipy.io
2
3
4 mat=scipy.io.loadmat("PeaksData.mat")
5
6 y=mat["Ct"]
7 x=mat["Yt"]
8
9 y_val=mat["Cv"]
10 x_val=mat["Yv"]
11
12
13 optimize_NN(x.T,y.T,x_val.T,y_val.T,depth=4,lr=0.1,batch_size=16,net_width=50,itors=150)
14 print("")

```


```

the amount of layers is 4
2 Final accuracy on train is 0.93476
Final accuracy on val is 0.93232

```



Loss on Peaks data set , using SGD



Experiments and failures network design :

We spent a lot of time trying to handle vanishing or exploding gradient, during the model arch process we had a hard time understanding why the model didn't learn. The problem after debugging was either saturation of the model due to tanh sending all the values to 1 or minus 1 because the weights are too big, or because ReLU exploded the gradient. We were unaware of the start that exploding gradient was the case and though perhaps the learning rate was too big, or maybe the batch-size. At start we tried to lower the size of the weights to enhance stability, but after a certain epoch, for each learning we choose, we either stopped progressing at all because he decreased too much, or the loss starting to increase and accuracy decreased.

Therefore we had to try different combination of tanh and relu in order to maintain stability, we tried to make tanh encapsulate the layer and ReLU serve as an activation layer for the middle step in the RNN network but the gradient still exploded, after debugging we realized which part of the network tends to explode and applied tanh there, which helped the stability of the network and finally learning could begin.

Hyper parameter searching Peaks data :

After finding a configuration which doesn't explode the gradient or vanishes it, we explore the hyper parameters of learning rate, batch-size and training time and depth.

Depth and epochs:

First, the width of the network is 50. We expect to see an increase in accuracy as the depth increases until a certain point, for a network with 2 layers we achieved 92.125 accuracy on train and 92.075 in convergence, which we reached after 50 iteration,

in a forward NN with 4 layers 150 iterations were enough for convergence , which peaked at 93.2% for train and 93% for test. we chose the depth to be 4 epochs amount to be 150.

Learning rate :

Until now we used a LR of 0.01 , this section will experiment with different learning rates.

LR=0.001

A smaller learning rate would converge more slowly but could increase the final accuracy , thus we experimented with 0.01 and 0.001 , first we checked the convergence point for 0.01 and then observed if 250 iters are enough to achieve convergence for a smaller learning rate , in the first attempt the accuracy plummeted to 40% in convergence , we assumed because we were stuck in a local minima and perhaps the initialization was unfavourable , thus we ran the learning procedure a few times and achieved the same results, therefore we understood that for a depth of 4 a lr smaller than 0.01 isn't beneficial to learning.

LR=0.1

The next logical step is observing higher learning rate , for a learning of 0.1 the learning procedure occurred as expected , and the achieved accuracy increased , after 150 epochs we reached 93.7% for train and 93.2% for the test. This learning rate seems to converge faster and attain higher acc so we chose him. Remark : We're using a decaying learning rate of $(0.99)^{epoch} * lr$

Batch Size :

Until now we used a batchsize of 16 , We'll attempt to see the accuracy for 8 , 16 , 32 batchsize , after 150 epochs using network width=50 , lr=0.1 , depth=4

Batch-size == 8 :

Final accuracy on train is 0.93496 Final accuracy on val is 0.93152 Compared to the previous batch_size of 16 of 93.7 and 93.2 , we assumed the training size is lower

Batch-size == 32

Final accuracy on train is 0.93548 Final accuracy on val is 0.93104 In both cases , the accuracy is lower , but the difference doesn't seem to be major between 8,16 and 32.

Therefore the final hyper-parameters for the network are : Width=50 lr=0.1 depth=4 batch-size=16 Epochs=150 and they correspond for 93.7% acc on train and 93.2% acc on test for Peaks dataset

Question 2 Part 4 :

Optimization on GMM datasets using ResNet

```
2 def init_weights_forward_ResNet(depth,WIDTH_NET,ORDER_NORM,NUM_CLASSES,INPUT_LEN):
3     """
4     Initialization of the weights for a resnet neural network
5     :param depth:
6     :param WIDTH_NET:
7     :param ORDER_NORM:
8     :param NUM_CLASSES:
9     :param INPUT_LEN:
10    :return:
11    """
12
13    # Init of the W matrices
14    w=[np.random.random((WIDTH_NET, INPUT_LEN))/ORDER_NORM,np.random.random((WIDTH_NET,
15        INPUT_LEN))/ORDER_NORM]
16    for i in range(depth-2):
17        w.append([np.random.random((WIDTH_NET, WIDTH_NET)) / ORDER_NORM,
18            np.random.random((WIDTH_NET, WIDTH_NET))/ORDER_NORM])
```

```

18 w.append([np.random.random((NUM_CLASSES,WIDTH_NET))/ORDER_NORM,np.random.random((NUM_CLASSES,W
20
21 # init of biases
22 b=[np.random.random(WIDTH_NET)/ORDER_NORM,np.random.random(WIDTH_NET)/ORDER_NORM]]
23 for i in range(depth-2):
24     b.append([np.random.random(WIDTH_NET)/ORDER_NORM,np.random.random(WIDTH_NET)/ORDER_NORM])
25 b.append([np.random.random(NUM_CLASSES) / ORDER_NORM,np.random.random(NUM_CLASSES) /
26     ORDER_NORM])
27 return w,b
28
29 def
30 optimize_RNN(x,y,x_val,y_val,lr=0.001,plot=True,batch_size=32,data_set="Peaks",depth=2,itera=5
31
32 #Flag for sizes in dependnece of the data
33 WIDTH_NET = width
34 if data_set=="Peaks":
35     ORDER_NORM = 1
36     NUM_CLASSES = 5
37     INPUT_LEN=2
38 if data_set=="GMM":
39     ORDER_NORM = 1
40     NUM_CLASSES = 5
41     INPUT_LEN=5
42
43 w,b=init_weights_forward_ResNet(depth=depth,WIDTH_NET=WIDTH_NET,ORDER_NORM=ORDER_NORM,NUM_CLAS
44 loss_train=[]
45 loss_val=[]
46 acc_train=[]
47 acc_val=[]
48 train_size = x.shape[0]
49
50 val_size=6250
51 #Each iteration is an epoch
52 for k in range(iters):
53     #Exponenet learning rate decay
54     curr_lr=np.power(0.999,k)*lr
55     for j in range(train_size // batch_size):
56         #Batch sampling and backwards
57         indices=np.random.randint(0, train_size, batch_size)
58         dw,db,dx=backward_rnn(x[indices,:],y[indices:],w,b)
59
60         for i in range(len(w)):
61             b[i][0]=b[i][0]-curr_lr*db[i][0]
62             b[i][1] = b[i][1] - curr_lr * db[i][1]
63             w[i][0]=w[i][0]-curr_lr*dw[i][0]
64             w[i][1] = w[i][1] - curr_lr * dw[i][1]
65
66
67 #calc accuracy on train
68 probs,hl,loss,_,_=RNN_forward(x[:, :], y[:, :], w, b, np.tanh)
69 agreements = np.argmax(probs, axis=1) == np.argmax(y[:, :], axis=1)
70 acc_train.append(agreements.sum() / train_size)
71 loss_train.append(loss)
72 #calc accuracy on validation

```

```

74     probs, hl, loss, _, _ = RNN_forward(x_val[:, :], y_val[:, :], w, b, np.tanh)
75     agreements = np.argmax(probs, axis=1) == np.argmax(y_val[:, :], axis=1)
76     acc_val.append(agreements.sum() / val_size)
77     loss_val.append(loss)
78     if plot:
79
80         plt.plot(acc_train, label="acc train")
81         plt.plot(acc_val, label="acc val")
82         plt.legend()
83         print(f'final accuracy for train is {acc_train[-1]}')
84         print(f'final accuracy for val is {acc_val[-1]}')
85
86         plt.title(f"Accuracy on {data_set} data set , using SGD")
87         plt.xlabel("Epochs")
88         plt.ylabel("Accuracy")
89         plt.show()
90
91         plt.plot(loss_train, label="loss train")
92         plt.plot(loss_val, label="loss val")
93         plt.legend()
94         plt.title(f"Loss on {data_set} data set , using SGD")
95         plt.yscale("log")
96         plt.xlabel("Epochs")
97         plt.ylabel("Loss")
98         plt.show()
99
100
101
102
103
104     return w, loss , acc_val, acc_train

```

```


mat=scipy.io.loadmat("GMMData.mat")
2
y=mat["Ct"]
4 x=mat["Yt"]
5
6
y_val=mat["Cv"]
8 x_val=mat["Yv"]
optimize_RNN(x.T,y.T,x_val.T,y_val.T,data_set="GMM",depth=3,width=20,itors=150)
10 print("")

```


```

final accuracy for train is 0.95344
2 final accuracy for val is 0.9536

```



Loss on GMM data set , using SGD



Experiments and failures network arch:

We spent a lot of time trying to handle vanishing or exploding gradient, during the model arch process we had a hard time understanding why the model didn't learn. The problem after debugging was either saturation of the model due to tanh sending all the values to 1 or minus 1 because the weights are too big, or because ReLU exploded the gradient. We were unaware of the start that exploding gradient was the case and though perhaps the learning rate was too big, or maybe the batchsize. At start we tried to lower the size of the weights to enhance stability, but after a certain epoch, for each learning we choose, we either stopped progressing at all because he decreased too much, or the loss starting to increase and accuracy decreased.

Therefore we had to try different combination of tanh and relu in order to maintain stability, we tried to make tanh encapsulate the layer and ReLU serve as an activation layer for the middle step in the RNN network but the gradient still exploded, after debugging we realized which part of the network tends to explode and applied tanh there, which helped the stability of the network and finally learning could begin.

Stability and depth :

As we increased the depth the stability decreased, which we balanced out using the width of the network. For a depth of 4 the hidden layers size had to be 10 for the training to remain stable, for a depth of 3 we had to use 20 layers. We believe we could improve the results if we managed to stabilize the training and we used those hyper-parameters for testing.

Depth = 2 network width =100 , iter =150

A shallow net results in : final accuracy for train is 0.94508 final accuracy for val is 0.94816

Depth =3 , network width=20 , iter =150

We can see that they still remain competitive even with a smaller amount of weights , and the best result were achieved using the current configuration , the amount of weights is still large , and the depth allow more non-linear patterns to emerge then a shallow network with more weights.

final accuracy for train is 0.95224 final accuracy for val is 0.9544

Depth =4 , network width=10 , iter =150

a much deeper net , with less weights overall despite being deeper managed to surpass the previous accuracy with a faster training time. We believe the results stem from the network ability to express more complicated patterns as the depth increases. final accuracy for train is 0.94668 final accuracy for val is 0.94704

Question 2.5 - only 200 samples


```
mat=scipy.io.loadmat("GMMData.mat")
2 import random

4 y=mat["Ct"]
  x=mat["Yt"]
6 indices=random.sample(range((x.shape[1])),200)

8 x=x[:,indices]

10 y=y[:,indices]
   y_val=mat["Cv"]
12 x_val=mat["Yv"]
   optimize_RNN(x.T,y.T,x_val.T,y_val.T,data_set="GMM",depth=2,width=100,itors=300,batch_size=8,lr=0.
14
   print("")

1 final accuracy for train is 0.995
  final accuracy for val is 0.91504
```



Loss on GMM data set , using SGD



Experiments

depth of 3 , width 20 , lr = 0.01 , batchsize =8

final accuracy for train is 0.99 final accuracy for val is 0.87968

depth of 4 , width 10 , lr = 0.01 , batchsize = 8

final accuracy for train is 0.985 final accuracy for val is 0.82288

depth of 2 , width 100 , lr = 0.01 , batchsize =8

final accuracy for train is 0.995 final accuracy for val is 0.91504

The optimal result is 1.0 since we expect extreme over fitting , we expect a network to over fit better as the amount of weights increases and therefor the result is optimal with width of 100 and depth of 2.

How did the result changes

The train accuracy increased dramatically to 0.995 , a perfect fitting of the data , which makes sense since the training data is 200 , compared to 25,000 therefor it is easier to over-fit with a smaller sample size. But the validation accuracy is lower compared to using the whole dataset , which indicate over-fitting. When we used the entire dataset the training accuracy and evaluation accuracy (and losses) remained similar , but a big disparity is present when using only 200 samples.